

The instruction set architecture (ISA)

The different aspects of ISA design

Introduction

From instruction to ISA

ISA Design

Conclusion

An ISA is an abstraction for computation

Software

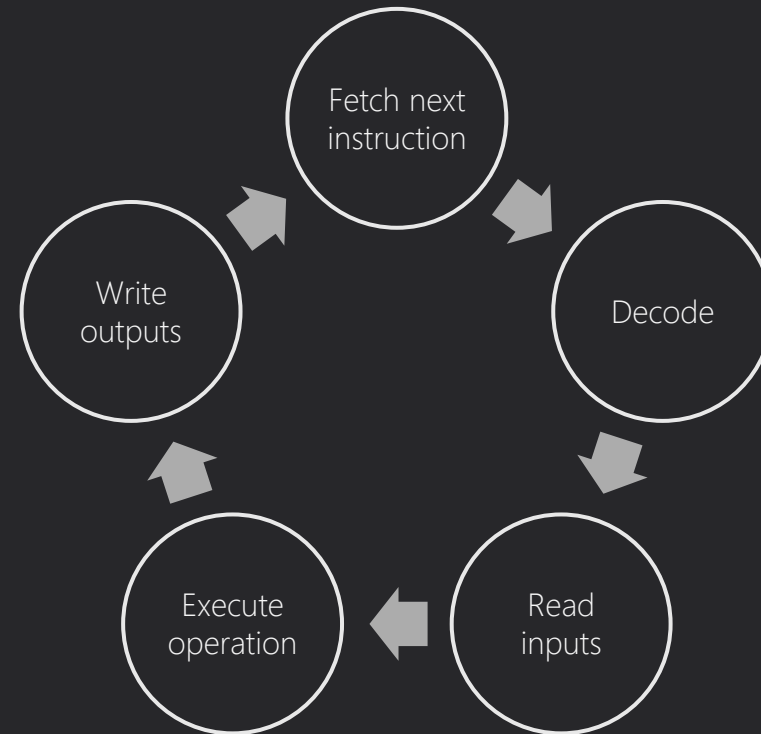
Instruction Set Architecture

Hardware

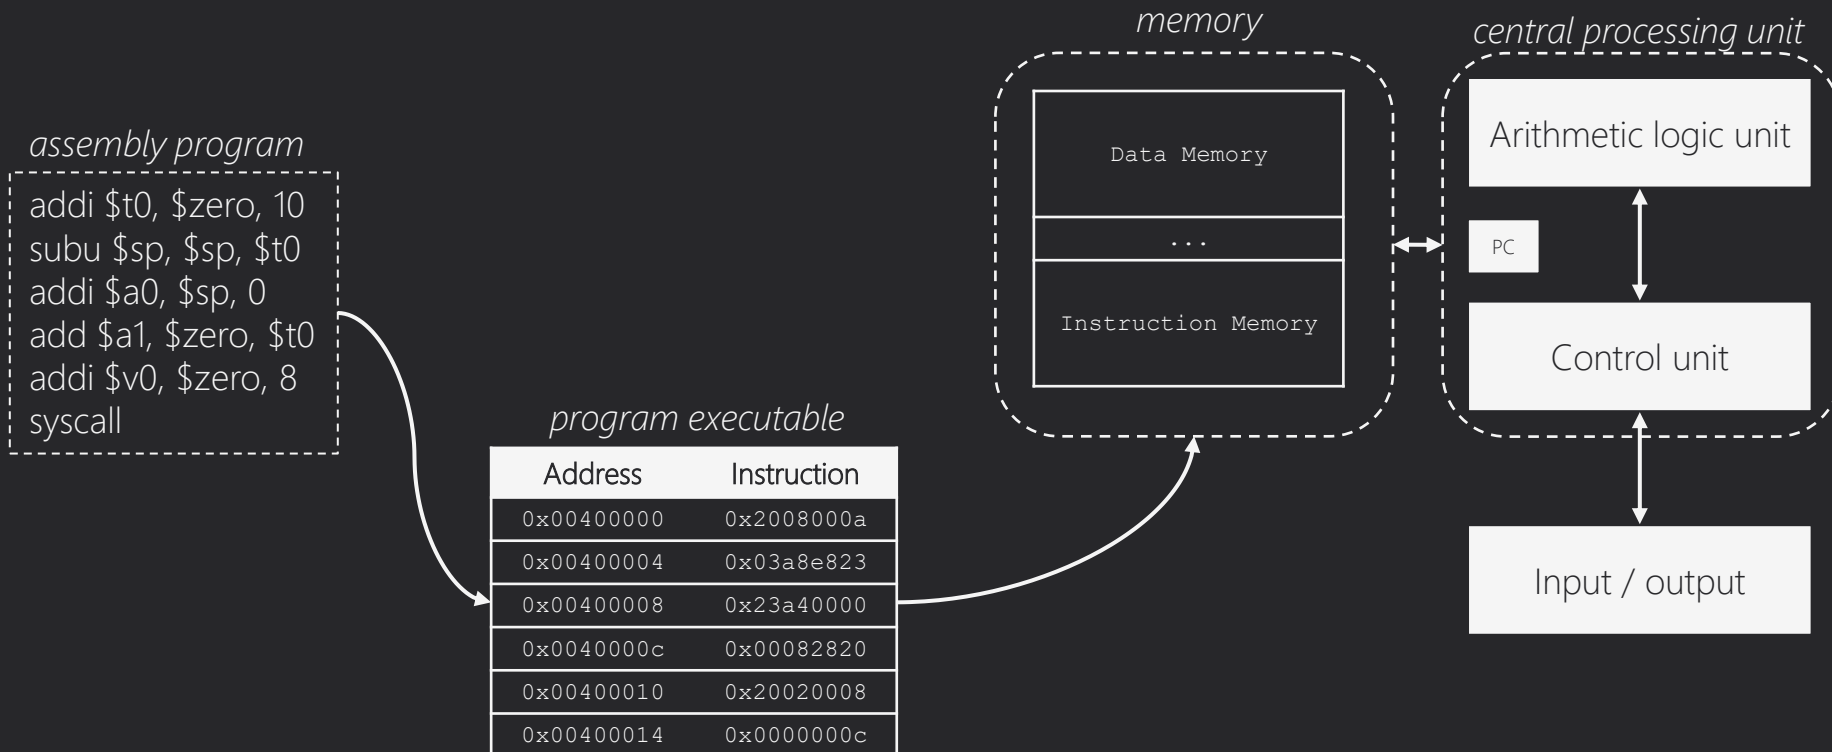
- Many different high-level languages
- Only a few major ISAs
- Many different processors on the market

The sequential state model

- The model of all major ISAs
- Processors repeat the loop (right)
- State updates are sequential
 - Instruction X finishes before $X + 1$
 - Matches a programmer's mental model of when memory is updated



The stored program concept



What's next

- An iterative build-up from instructions to instruction set to ISA
 - The different aspects of an ISA
- A brief comparison and look at some design decisions from existing instruction set architectures

The instruction

The instruction set

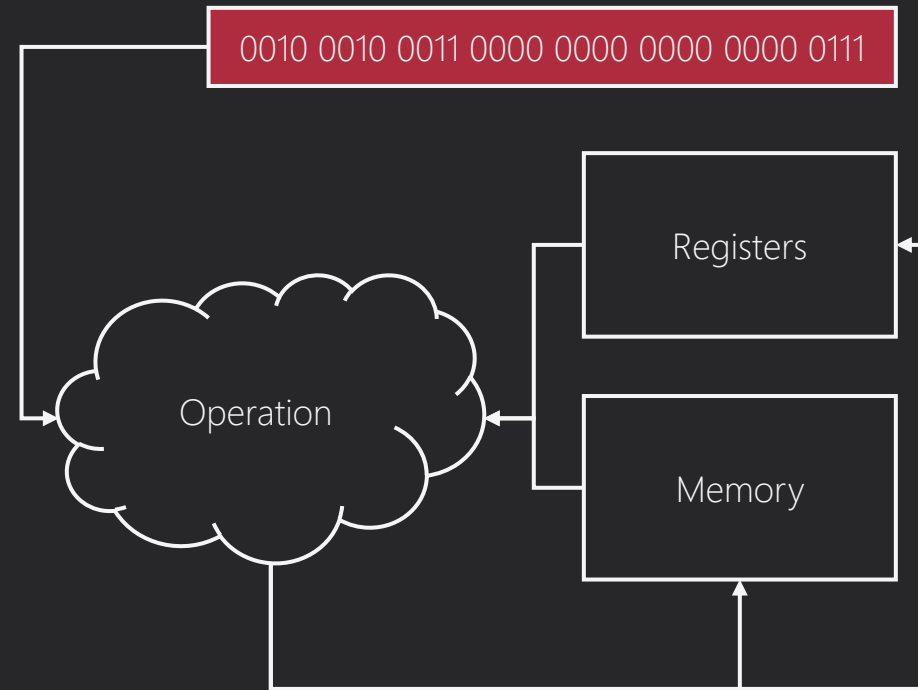
The instruction set architecture

From instruction to ISA

Some of the easier-to-see parts of an ISA

The instruction

- An *instruction* defines the
 - operation performed by the hardware
 - operands to use for that operation
- Operations mutate state
- An *operand* is an input to the operation
 - Can be a register, memory location, or immediate
 - An immediate value is data encoded into the instruction itself



The instruction set

- Simply: a set of instructions
- The instructions can be grouped by the types of operation they perform
- A subset of the instruction set is the “common case”
 - i.e., the subset makes up most of the instructions executed by software

Operation type	MIPS examples
Arithmetic	<code>add, addi, addu</code>
Logical	<code>and, or</code>
Control	<code>beq, jal</code>
Data transfer	<code>lw, sw, lui</code>
System	<code>syscall</code>
Floating point	<code>add.s, div.s</code>

Operand sizes (and types)

Common to most ISAs

- Sizes
 - Byte (8 bits)
 - Half word (16 bits)
 - Word (32 bits)
 - Double word (64 bits)
- Types
 - Single-precision floating point (32 bits)
 - Double-precision floating point (64 bits)
 - Signed integers use 2's complement

Not-so-common examples

- Character strings
 - Each character is 1 byte (not so useful with Unicode)
- “Decimal” formats
 - Because not all fractions can be represented exactly in floating point

The instruction set architecture

Aspects

- Operations ([earlier slide](#))
- Data types ([earlier slide](#))
- Instruction format
- Operand model
- Addressing
- Control

Design philosophies

- *Reduced Instruction Set Computer* (RISC)
 - A “small” set of instructions
 - Instructions perform “simple” operations
 - A focus on the common case
- *Complex Instruction Set Computer* (CISC)
 - A large set of instructions
 - Instructions perform both simple and complex (for hardware) operations
 - Support for the uncommon case

Instruction format

Length and encoding

- Fixed-length
 - e.g., MIPS: 4 bytes
 - Three types of encodings (R-, I-, and J-type)
- Variable-length
 - e.g., x86: 1 to 15 bytes
- Hybrid
 - e.g., Arm's Thumb2: 2 or 4 bytes

Impact

- Program size
 - Important in embedded systems
 - Fixed-length does poorly here
- Hardware complexity
 - Hardware must decode instructions before it can execute them
 - Variable-length does poorly here

MIPS instruction formats

- Three formats
 - R-type (operations only use registers)
 - I-type (operations use an immediate)
 - J-type (for “jump” operations)
- Notice the regularity
 - op for all formats
 - rs and rt for R- and I-type


R-type



I-type

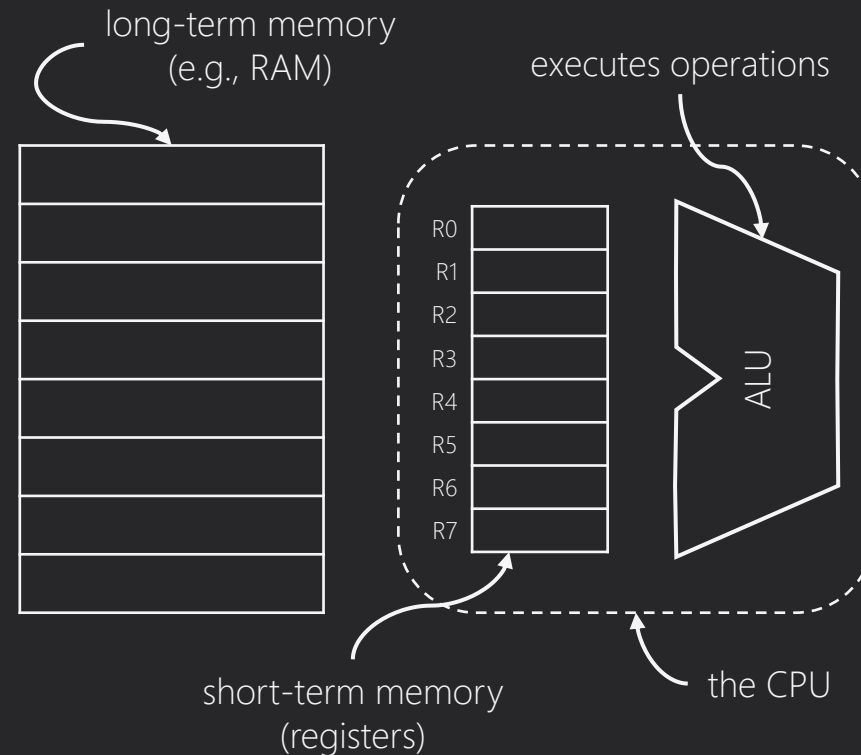


J-type



Classifying ISAs by operand model

- Modern ISAs define a set of general-purpose registers for “internal storage”
- A *register-memory architecture* has instructions that operate on memory or registers
 - e.g., an add operation may have one operand in memory and another in a register
- A *register-register architecture* only operates on registers
 - Data in memory must be transferred to/from registers with dedicated instructions
 - Also called a *load-store architecture*



The set of general-purpose registers

- MIPS defines thirty-two 32-bit registers
 - With conventions on how to use those registers
 - A register can be uniquely identified with 5 bits ($2^5 = 32$)
- The 32-bit version of the x86 architecture (IA-32) defines eight 32-bit general-purpose registers

Name	Address	Description
\$0	0	The value 0
\$at	1	A temporary value for the "assembler"
\$v0-\$v1	2-3	Return values for procedure calls
\$a0-\$a3	4-7	Arguments for procedure calls
\$t0-\$t7	8-15	Temporary variables
\$s0-\$s7	16-23	"Saved" variables
\$t8-\$t9	24-25	More temporary variables
\$k0-\$k1	26-27	Temporary variables for operating systems
\$gp	28	The global pointer
\$sp	29	The stack pointer
\$fp	30	The frame pointer
\$ra	31	The return address for procedure calls

Interpreting memory addresses

- Modern ISAs use byte-addressable memory
 - 32-bit addresses (old), 64-bit addresses (today)
- Byte-addressable memories can be *big-endian* or *little-endian*
 - i.e., byte 0 of a word contains different data on a big-endian and little-endian machine
- Memory accesses to data larger than 1 byte must (typically) be aligned
 - e.g., accesses to words (4 bytes) must use (or are preferred to use) addresses that are a multiple of 4

	Big-Endian				Little-Endian			
Byte address	0	1	2	3	3	2	1	0
Data	12	34	56	78	12	34	56	78
	MSB		LSB		MSB		LSB	

Word-sized data accesses	0	1	2	3	4	5	6	7
Addresses 0 and 4	Aligned				Aligned			
Addresses 1 and 5		Misaligned				Misaligned		
Addresses 2 and 6			Misaligned				Misaligned	
Addresses 3 and 7				Misaligned				M...

Addressing modes

- An *addressing mode* defines how operands are identified
- Addressing modes specify constants, registers, and locations in memory

MIPS addressing mode	The instruction specifies...	Example
Register-only	Only register addresses.	$[rd] = [rs] + [rt]$
Immediate	A value (not an address).	$[rt] = [rs] + \text{immediate}$
Base-displacement	A displacement from an address that is in a register.	$\text{address} = [rs] + \text{immediate}$ $[rt] = \text{mem}[\text{address}]$
PC-relative	A value relative to the address in the PC.	$\text{address} = [\text{PC}] + \text{imm}$

Conditional execution (control)

Testing for conditions

- Compare-and-branch
 - e.g., MIPS' branch-if-equal (`beq`)
 - e.g., MIPS' branch-if-greater-than-zero (`bgtz`)
- Condition register
 - e.g., MIPS' set-less-than (`slt`) instruction
 - Must then use a branch instruction like `beq` or `bne`
- Condition code
 - A condition code register is set based on a previous operation
 - Branch based on value of condition code

Computing target addresses

- Need to calculate the target address to store in the PC
- Example addressing modes
 - PC-relative
 - Direct (or pseudo-direct)
 - Indirect

ISA Design



The driving forces behind the design of ISAs

Well-designed ISAs

A brief comparison


x86 and implementation

Aspects of a well-designed ISA



Software: easy to express
computation

Compiler complexity



Hardware: easy to implement

Trade-offs: performance, power, cost



Future-proof: easy to maintain
compatibility

Trends in software

Trends in technology

A brief ISA comparison

- Like software, ISAs have “versions”
 - e.g., Armv8-R is a version of Arm64

Name	Operand model	Encoding	Control	General-purpose registers
MIPS32	register-register	• Fixed (32-bit)	• Compare and branch	32 32-bit
x86 (64-bit)	register-memory	• Variable (1 to 15 bytes)	• Condition code	16 64-bit
Arm32	register-register	• Fixed (32-bit)	• Condition code • Compare-and-branch	15 32-bit
Arm64	register-register	• Hybrid (32- or 64-bit)	• Condition code • Compare-and-branch	31 64-bit
RISC-V	register-register	• Fixed (base) • Variable (extensions)	• Compare-and-branch	32 64-bit

Simplifying implementation through translation

Problem and solution

- Due to its history, x86 has a set of difficult-to-implement instructions
 - Can't remove (compatibility)
- Intel's solution? "Complex" decode hardware to simplify the rest of the implementation

Example

- Original instruction:

```
push $eax           # push [eax] onto the stack
```

- Hardware translation (after decode):

```
addi $esp, $esp, -4    # make space on the stack  
sw $eax, 0($esp)       # put [eax] in allocated space
```

How different are modern ISAs?


Commonalities

- Data type support
- General-purpose registers
- Support for frequently-used operations (e.g., add, multiply)
- Memory alignment restrictions

Differences

- Instruction format
- Design philosophies and historical inertia

Conclusion



Final questions and thoughts

Why define an instruction set architecture?

- One architecture has *many* different hardware implementations
 - e.g., Intel vs. AMD processors
 - e.g., Apple's M1 vs. M2 processors
- Software compiled for one architecture can run on many different processors
- Wait... What about the operating system?

Are ISAs only implemented in hardware?

- No, an ISA is just an interface
 - Can be implemented in software
 - Can be implemented in a mixture of hardware and software
- Many emulators (software) exist, especially for old architectures