



# Pipeline dimensions

---

The depth and width of a pipeline

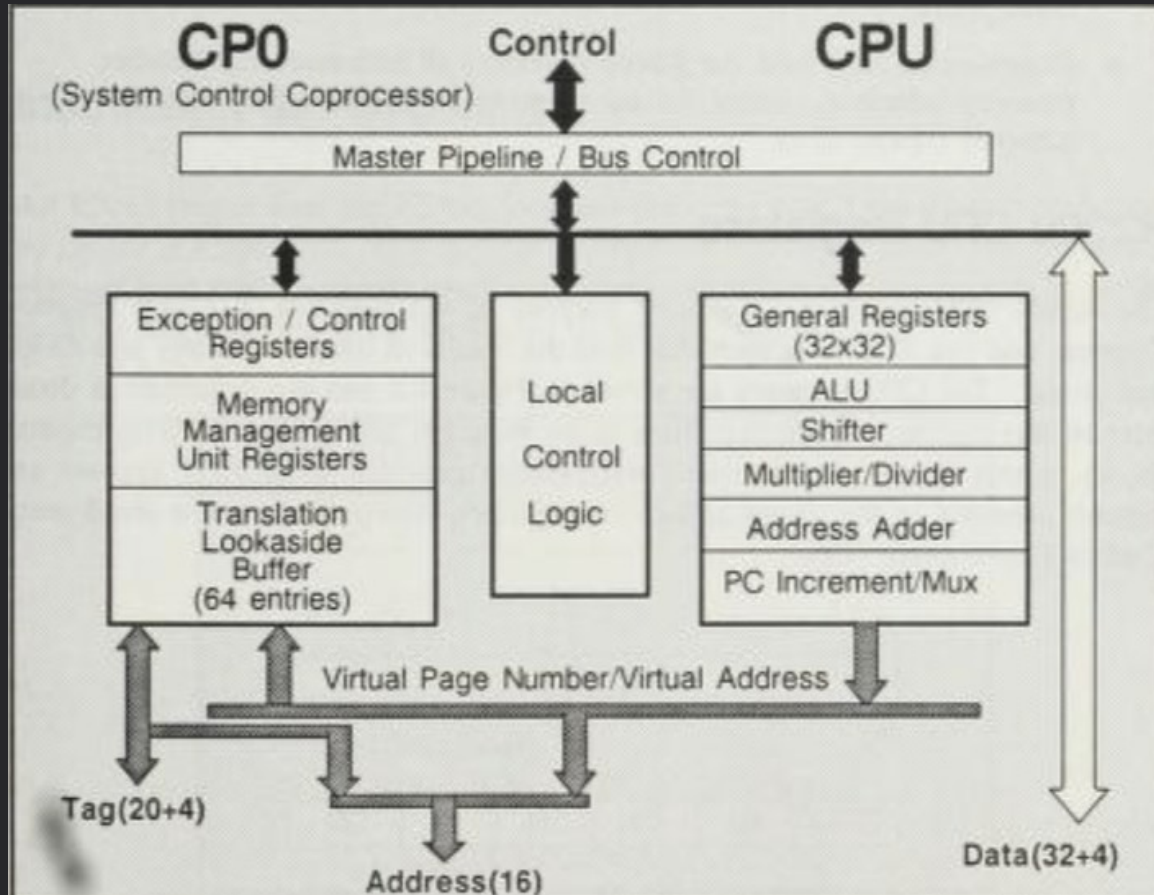
Introduction

Review and history

Improving performance

Conclusion

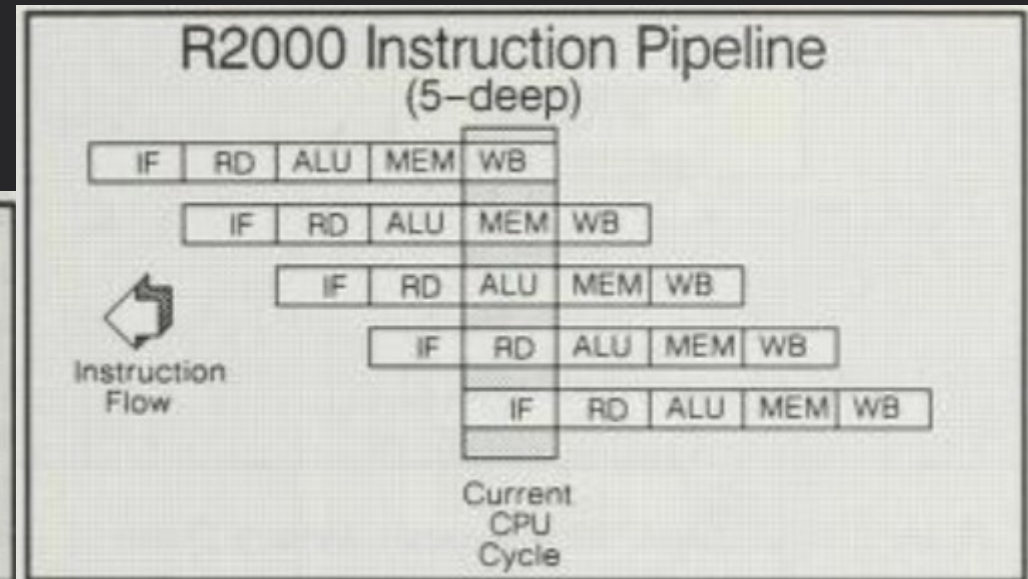
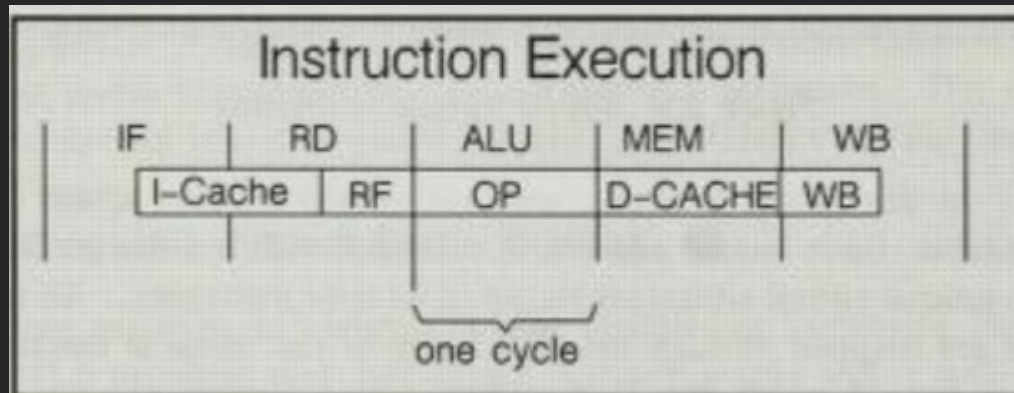
# The MIPS R2000 (1986)



Kane, Gerry. [MIPS RISC Architecture](#). Prentice-Hall, Inc., 1988.

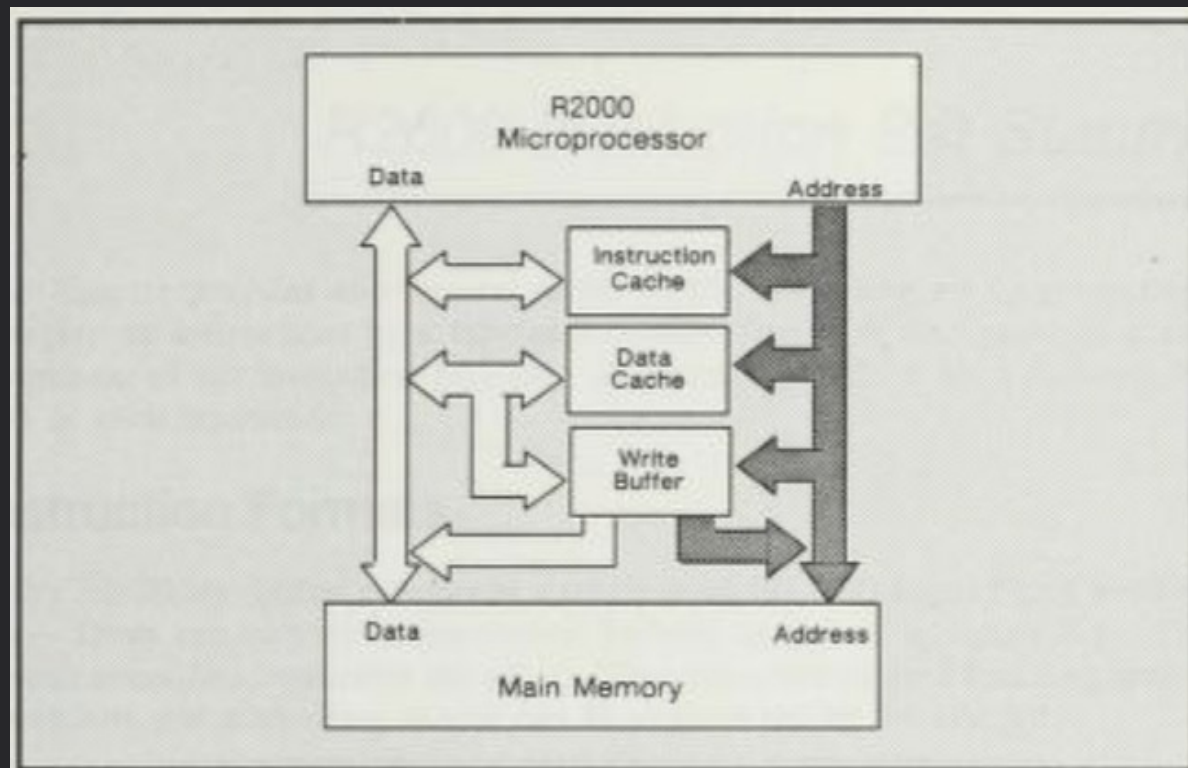
- Coprocessor
  - Generates addresses
  - Handles all control
  - Handles exceptions
- CPU
  - Pipelined processor
  - Autonomous multiplier/divider

# The MIPS R2000 pipeline



Kane, Gerry. [MIPS RISC Architecture](#). Prentice-Hall, Inc., 1988.

# The MIPS R2000 memory hierarchy



Kane, Gerry. [MIPS RISC Architecture](#). Prentice-Hall, Inc., 1988.

- The caches were actually “off chip”
  - But still closer and faster to access than main memory
- Used a write-through policy to keep main memory up-to-date

# How can we improve execution time?

$$\text{Execution Time} = \left( \frac{\text{instructions}}{\text{program}} \right) \left( \frac{\text{cycles}}{\text{instruction}} \right) \left( \frac{\text{seconds}}{\text{cycle}} \right)$$

- Reduce cycle time
  - Temporal parallelism
  - Pipeline depth
- Reduce CPI (increase IPC)
  - Spatial parallelism
  - Pipeline width

# Increasing pipeline depth

## Advantages

---

- Increases the frequency
  - Equivalently: Reduces the cycle time ( $T_C$ )


## Disadvantages

---

- Limited by pipeline registers
  - Sequencing overhead
  - Extra hardware
- Limited by hazards
  - More complex control and hazard detection
  - More hardware for bypassing
  - Branches resolved later in the pipeline; Penalty of a flush increases

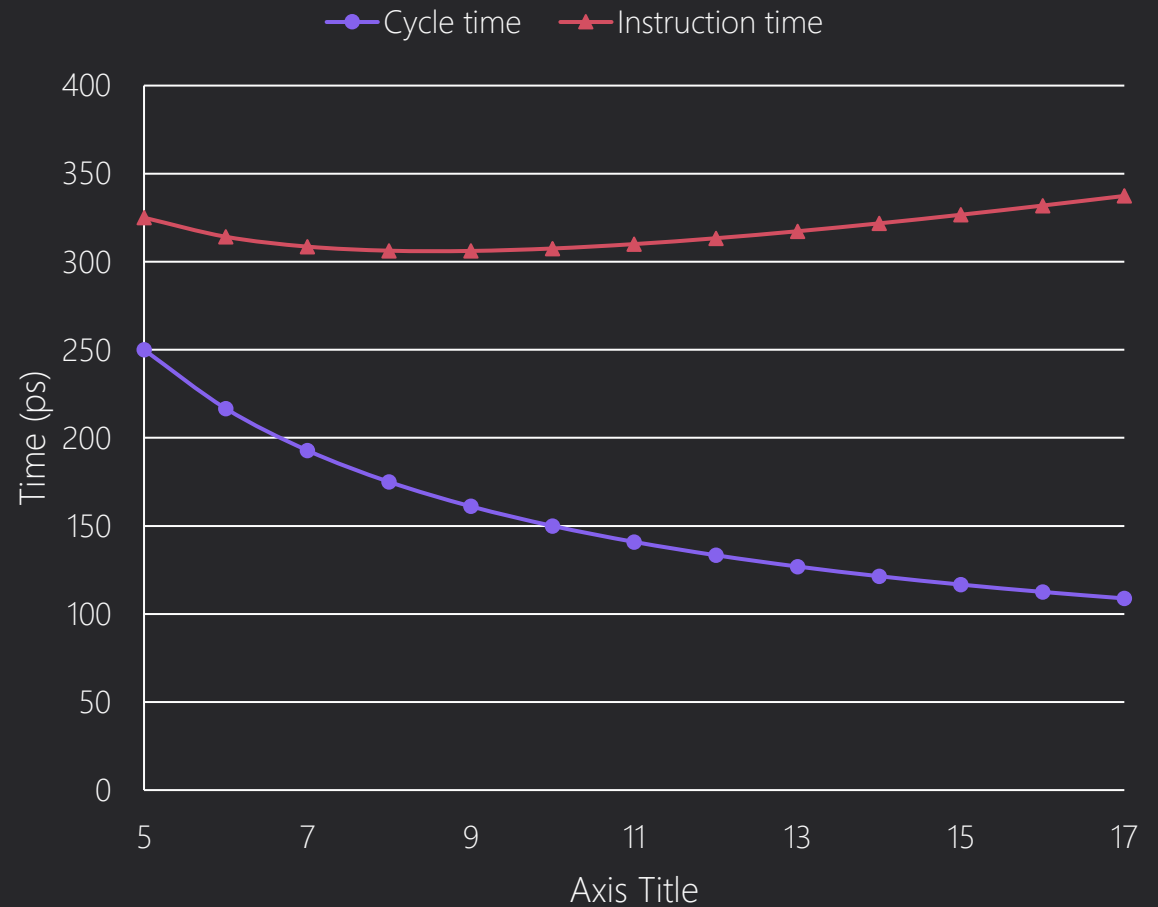
# Superpipelined CPU

When a CPU has more than the classical 5 stages, we say it is *superpipelined*.



# The limits of pipelining

- Suppose a single-cycle implementation can be clocked at 1 GHz (1000 ps cycle time)
- For each pipeline stage added:
  - Sequencing overhead of 50 ps
  - $N = 5$ ;  $CPI_{base} = 1.3$
  - $N > 5$ ;  $CPI_{penalty} = 0.15 * N$
- How does pipelining the datapath to  $N$  stages impact:
  - The cycle time ( $T_C$ )
  - The instruction time ( $CPI \times T_C$ )





# Increasing pipeline width

## Advantages

---

- Increases the *ideal IPC* beyond 1
  - Allows more than one instruction in a stage at a time
  - More options to exploit instruction-level parallelism

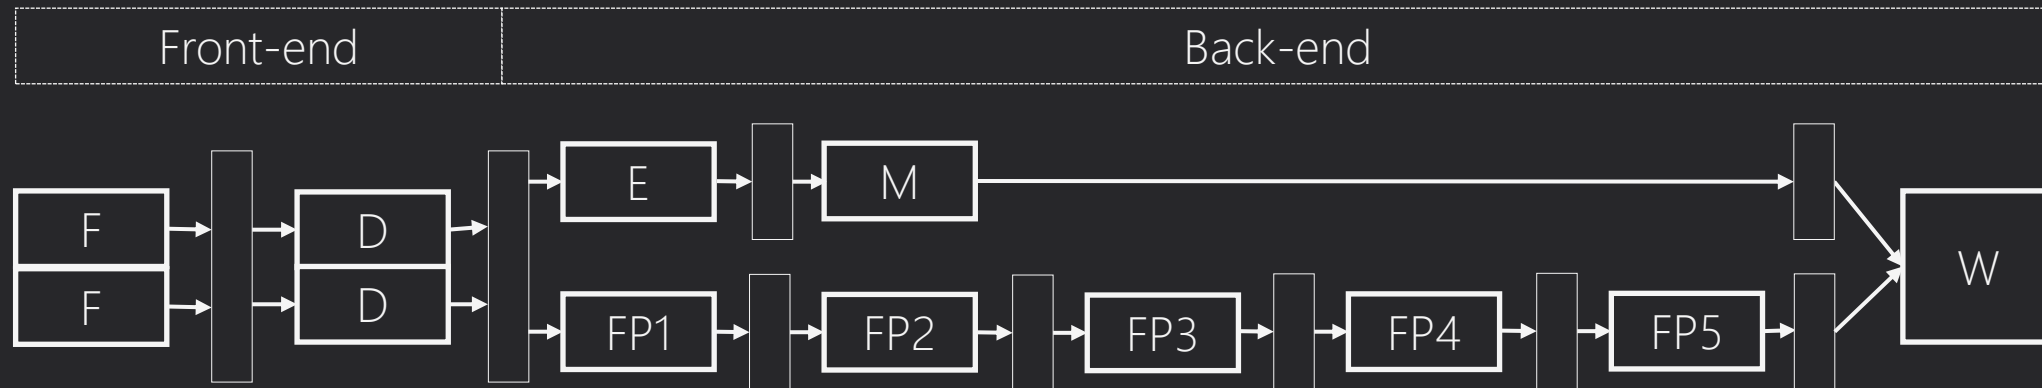
## Disadvantages

---

- Significant changes to our existing pipeline structure
- Complex scheduling
- Increased cost
  - Area
  - Power

# Superscalar CPU

- $m$  is the (ideal) number of instructions fetched into the pipeline
- If  $m = 1$ , the processor is simply *scalar*
- $m$ -way superscalar; here  $m = 2$



# The limits of superscalar

- Suppose an  $m$ -way superscalar with a frequency  $k$  times that of a scalar processor
- Theoretical speedup:  $m \times k$
- Large  $m$  needs lots of ILP
  - Limits to hardware discovering ILP
  - Limits to ILP
- Back-end needs at least  $m$  functional units
  - Number of forwarding paths grows quadratically with number of functional units
- More ports are costly
  - e.g., register file needs more read ports
- Power
  - Increasing  $m$  and  $k$  makes for a very power-hungry processor

# Conclusion



Recap of important points

# Improving execution time

Pipeline Dimension	Improvement	Trade-off
Depth	Better frequency (lower cycle time)	Diminishing returns due to: <ul style="list-style-type: none"><li>• Sequencing overhead</li><li>• Hazards</li></ul>
Width	Better ideal IPC ( $IPC_{ideal} = m$ )	Diminishing returns due to: <ul style="list-style-type: none"><li>• Limits of ILP and scheduling</li><li>• Hardware cost</li></ul>