

# Speculation in memory

---

Hiding memory latency from the  
program

Introduction

Prefetching overview

Types of prefetchers

Conclusion

# When is data brought into the cache?

## On-demand

---

- A cache block is only filled when a miss for that block occurs
- In an L1 cache, misses originate from load and store instructions

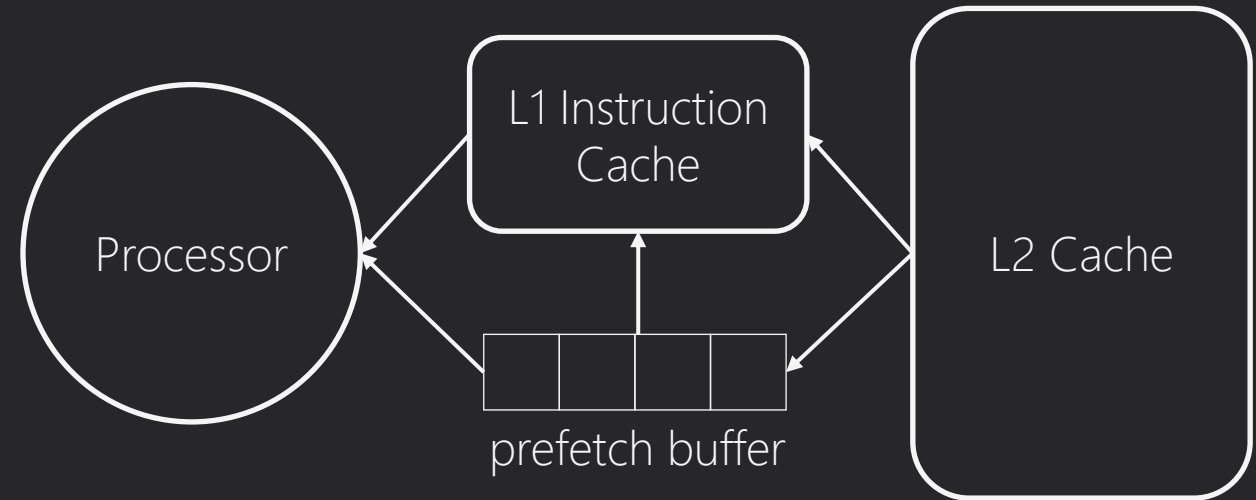
## Prefetch

---

- A cache block may be filled due to speculation
  - i.e., predict that this cache block is needed in the "near future"

# What is prefetching?

- A predictive technique that brings data into a cache before it is needed
- Prefetching can be done...
  - For instructions or data (or both)
  - Between any two levels of the memory hierarchy
- Goal: Increase the hit ratio of the cache



# What if the prefetcher mispredicted?

- Mispredictions do not impact program correctness
- Mispredictions may result in worse performance
  - e.g., the prefetch evicted a useful cache block
- Can prefetchers adapt to mispredictions?
  - Possible, but feedback mechanism is difficult and lengthy

# How are prefetchers assessed?

- $Accuracy = \frac{\text{useful prefetches}}{\text{number of prefetches}}$
- $Coverage = \frac{\text{useful prefetches}}{\text{number of misses (without prefetcher)}}$
- Timeliness
  - Too early: evicted useful data from the cache
  - Too late: a demand access had to wait (though the prefetch reduces the latency)
  - Too early (worst-case): two evictions (one for useful data, one for prefetched data)

# What are the disadvantages of prefetching?

- Prefetchers compete for hardware used by regular memory operations
- e.g., should not prefetch lines that are already in the cache
  - Contention: cache tags
- Cache pollution
  - e.g., evictions and “Too early (worst-case)” behaviour
- Bandwidth pollution
  - e.g., need to send requests to next level of the hierarchy (via a bus). Priority to the bus must be given to load and store instructions

# What is a sequential prefetcher?

- Prefetch an adjacent cache block to  $i$ 
  - Forward:  $i + 1$
  - Backward:  $i - 1$
- When to prefetch?
  - Always prefetch
  - Prefetch on miss
  - Other?
- What is the prefetch degree?
  - The number of blocks to prefetch  $N$



Fetch  $i$

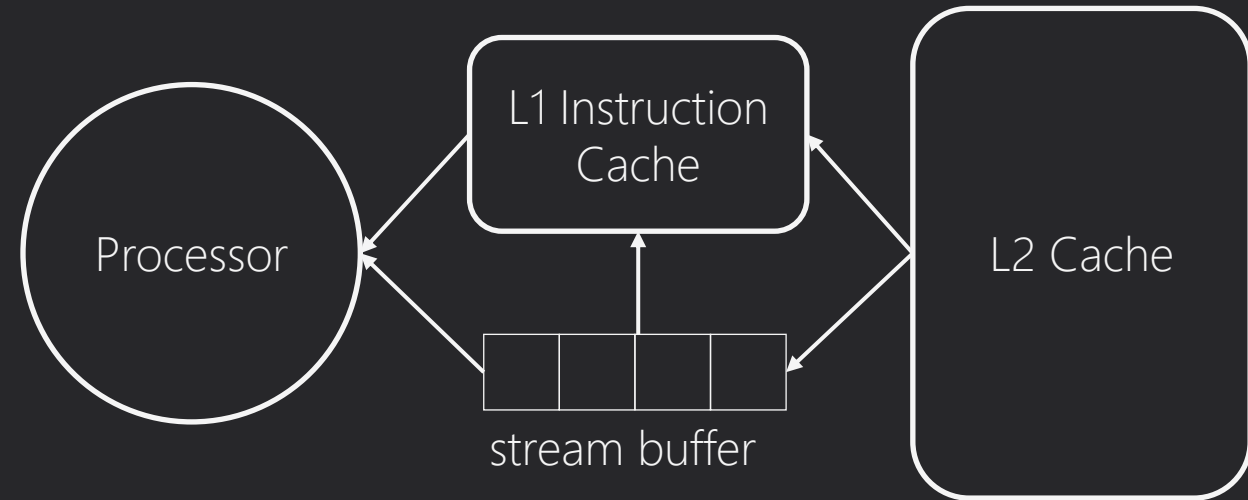


Prefetch  $i + 1$

| Address | Instruction       |
|---------|-------------------|
| 500     | LW r4, 0(r2)      |
| 504     | LW r5, 0(r3)      |
| 508     | MUL r6, r5, r4    |
| 512     | LW r7, 0(r1)      |
| 516     | ADD r7, r7, r6    |
| 520     | SW r7, 0(r1)      |
| 524     | ADD r2, r2, 4     |
| 528     | ADD r3, r3, 400   |
| 532     | ADD r11, r11, 1   |
| 536     | BNE r11, r13, 500 |

# Example: Alpha AXP 21064

- Add a *stream buffer* to hold prefetched data
- On access,
  - Check cache *and* stream buffer
  - If found in stream buffer, move to cache
- On miss,
  - Fetch missed block  $i$  into cache
  - Prefetch block  $i+1$  into stream buffer iff it is not already in the cache





# Analysis of sequential prefetching

## Advantages

---

- Adapts to different programs (and phases of programs)
  - Hence, different from just increasing block size
- Stream buffer avoids cache pollution
  - Only move data to cache when it is accessed
- Great for instruction caches (sequential access pattern)

## Disadvantages

---

- Timeliness
- Memory accesses patterns are not always sequential
- Stream buffers have some overhead
  - For instruction cache, one stream buffer is usually sufficient
- Poor for data caches

# Effective hardware prefetching

■ *Chen, Tien-Fu, and Jean-Loup Baer. "Effective hardware-based data prefetching for high-performance processors." IEEE transactions on computers 44.5 (1995): 609-623.*

# How do we classify memory access patterns?

```
int A[100][100];
int B[100][100];
int C[100][100];

for i = 1 to 100:
  for j = 1 to 100:
    for k = 1 to 100:
      A[i][j] += B[i][k] * C[k][j]
```

- Scalar
  - e.g., reference to  $i$
- Zero stride
  - e.g.,  $A[i][j]$  is the same in iterations of the inner-most loop
- Constant stride
  - e.g.,  $B[i][k]$  in the inner loop is accessing  $B[i][0], B[i][1], \dots, B[i][100]$
- Irregular
  - (not shown)
  - e.g., for linked lists

# Example: matrix multiply's inner loop


```
int A[100][100];
int B[100][100];
int C[100][100];

for i = 1 to 100:
  for j = 1 to 100:
    for k = 1 to 100:
      A[i][j] += B[i][k] * C[k][j]
```

| <i>Address</i> | <i>Instruction</i> | <i>Comment</i> | <i>Stride</i> |
|----------------|--------------------|----------------|---------------|
| 500            | LW r4, 0(r2)       | B[i,k]         | 4B            |
| 504            | LW r5, 0(r3)       | C[k,j]         | 400B          |
| 508            | MUL r6, r5, r4     |                |               |
| 512            | LW r7, 0(r1)       | A[i,j]         | 0B            |
| 516            | ADD r7, r7, r6     |                |               |
| 520            | SW r7, 0(r1)       |                |               |
| 524            | ADD r2, r2, 4      |                |               |
| 528            | ADD r3, r3, 400    |                |               |
| 532            | ADD r11, r11, 1    |                |               |
| 536            | BNE r11, r13, 500  |                |               |

# The reference prediction table (RPT)


- The tag: Address of the load or store instruction
- The prev\_addr: The last *effective address*
- The stride: The difference between the last two addresses
- The state: Whether we should make a prediction
- Mechanism
  1. Record the effective address
  2. Compute the stride
  3. Set the state (compare previous stride with the one just computed)
  4. Predict if the state is "steady"



# State of RPT rows

- Initial
  - e.g., when row is first accessed
- Transient
  - e.g., predictions are not yet stable
- Steady
  - e.g., predictions should be stable
- No prediction
  - e.g., to disable prefetching

Strides are updated on incorrect transitions.



# Example: matrix multiply and RPT (0)

| Address | Instruction       | Comment | Stride |
|---------|-------------------|---------|--------|
| 500     | LW r4, 0(r2)      | B[i,k]  | 4B     |
| 504     | LW r5, 0(r3)      | C[k,j]  | 400B   |
| 508     | MUL r6, r5, r4    |         |        |
| 512     | LW r7, 0(r1)      | A[i,j]  | 0B     |
| 516     | ADD r7, r7, r6    |         |        |
| 520     | SW r7, 0(r1)      |         |        |
| 524     | ADD r2, r2, 4     |         |        |
| 528     | ADD r3, r3, 400   |         |        |
| 532     | ADD r11, r11, 1   |         |        |
| 536     | BNE r11, r13, 500 |         |        |

| Initially empty |           |        |       |
|-----------------|-----------|--------|-------|
| tag             | prev_addr | stride | state |
|                 |           |        |       |
|                 |           |        |       |
|                 |           |        |       |
|                 |           |        |       |

# Example: matrix multiply and RPT (1)

| Address | Instruction       | Comment | Stride |
|---------|-------------------|---------|--------|
| 500     | LW r4, 0(r2)      | B[i,k]  | 4B     |
| 504     | LW r5, 0(r3)      | C[k,j]  | 400B   |
| 508     | MUL r6, r5, r4    |         |        |
| 512     | LW r7, 0(r1)      | A[i,j]  | 0B     |
| 516     | ADD r7, r7, r6    |         |        |
| 520     | SW r7, 0(r1)      |         |        |
| 524     | ADD r2, r2, 4     |         |        |
| 528     | ADD r3, r3, 400   |         |        |
| 532     | ADD r11, r11, 1   |         |        |
| 536     | BNE r11, r13, 500 |         |        |

*After iteration 1*

| tag | prev_addr | stride | state |
|-----|-----------|--------|-------|
| 500 | 50,000    | 0      | init  |
| 504 | 90,000    | 0      | init  |
|     |           |        |       |
| 512 | 10,000    | 0      | init  |



# Example: matrix multiply and RPT (2)

| Address | Instruction       | Comment | Stride |
|---------|-------------------|---------|--------|
| 500     | LW r4, 0(r2)      | B[i,k]  | 4B     |
| 504     | LW r5, 0(r3)      | C[k,j]  | 400B   |
| 508     | MUL r6, r5, r4    |         |        |
| 512     | LW r7, 0(r1)      | A[i,j]  | 0B     |
| 516     | ADD r7, r7, r6    |         |        |
| 520     | SW r7, 0(r1)      |         |        |
| 524     | ADD r2, r2, 4     |         |        |
| 528     | ADD r3, r3, 400   |         |        |
| 532     | ADD r11, r11, 1   |         |        |
| 536     | BNE r11, r13, 500 |         |        |

*After iteration 2*

| tag | prev_addr | stride | state     |
|-----|-----------|--------|-----------|
| 500 | 50,004    | 4      | transient |
| 504 | 90,400    | 400    | transient |
|     |           |        |           |
| 512 | 10,000    | 0      | steady    |

# Example: matrix multiply and RPT (3)

| Address | Instruction       | Comment | Stride |
|---------|-------------------|---------|--------|
| 500     | LW r4, 0(r2)      | B[i,k]  | 4B     |
| 504     | LW r5, 0(r3)      | C[k,j]  | 400B   |
| 508     | MUL r6, r5, r4    |         |        |
| 512     | LW r7, 0(r1)      | A[i,j]  | 0B     |
| 516     | ADD r7, r7, r6    |         |        |
| 520     | SW r7, 0(r1)      |         |        |
| 524     | ADD r2, r2, 4     |         |        |
| 528     | ADD r3, r3, 400   |         |        |
| 532     | ADD r11, r11, 1   |         |        |
| 536     | BNE r11, r13, 500 |         |        |

*After iteration 3*

| tag | prev_addr | stride | state  |
|-----|-----------|--------|--------|
| 500 | 50,008    | 4      | steady |
| 504 | 90,800    | 400    | steady |
|     |           |        |        |
| 512 | 10,000    | 0      | steady |

# Analysis of RPT

- Different load and store instructions are separated via their PC
  - Like a cache
- Scalar and zero stride need one transition to steady state
  - From init to steady
- Constant stride need two transitions to steady state
  - From init to transient to steady
- Irregular access patterns don't reach steady state
  - Avoid cache pollution

TABLE II  
CHARACTERISTICS OF BENCHMARKS

| Name     | data references        |                    |           | branch<br>pred.<br>miss ratio |
|----------|------------------------|--------------------|-----------|-------------------------------|
|          | scalar,<br>zero stride | constant<br>stride | irregular |                               |
| Tomcatv  | 0.312                  | 0.682              | 0.006     | 0.005                         |
| Fpppp    | 0.981                  | 0.006              | 0.014     | 0.110                         |
| Matrix   | 0.059                  | 0.921              | 0.021     | 0.073                         |
| Spice    | 0.581                  | 0.239              | 0.180     | 0.060                         |
| Doduc    | 0.692                  | 0.154              | 0.154     | 0.120                         |
| Nasa     | 0.006                  | 0.989              | 0.003     | 0.008                         |
| Eqntott  | 0.338                  | 0.574              | 0.088     | 0.069                         |
| Espresso | 0.460                  | 0.424              | 0.116     | 0.055                         |
| Gcc      | 0.516                  | 0.120              | 0.365     | 0.204                         |
| Xlisp    | 0.440                  | 0.078              | 0.482     | 0.156                         |

# Conclusion

- Most modern processors use sequential prefetching for their instruction cache
- Prefetchers can be complex
  - Example shown is stride prefetching
  - Other examples exist (e.g., for pointer-based structures)
- Prefetching can be done in software, too
  - e.g., explicit instructions to initiate a prefetch
  - Programmer needs to ensure prefetches are timely